

HM10_final

Sedreh

5/29/2019

```
library("ggplot2")
library(data.table)
library(dplyr)
library(GGally)
library(DataExplorer)
library(plyr)
```

```
#Anscombe's quartet comprises four datasets
```

```
anscombe <- readRDS("/home/sedreh/ITMO/semester2/Statistic-R/10/anscombe.rds")
head(anscombe)
```

```
##      x      y set
## 1 10 8.04    1
## 2  8 6.95    1
## 3 13 7.58    1
## 4  9 8.81    1
## 5 11 8.33    1
## 6 14 9.96    1
```


```
str(anscombe)
```

```
## 'data.frame':   44 obs. of  3 variables:
## $ x : num  10 8 13 9 11 14 6 4 12 7 ...
## $ y : num  8.04 6.95 7.58 8.81 8.33 ...
## $ set: num  1 1 1 1 1 1 1 1 1 1 ...
```

```
#Scatter plot facettted by set
```

```
p <- ggplot(anscombe, aes(x, y, color = as.factor(set))) +
  geom_point()+
  facet_wrap(~set)+
  ggtitle("Scatter plot for anscombe dataset ")+
  theme_classic()
p
```

Scatter plot for anscombe dataset



#1: possitive linear relationship

#2: the relationship is non_linear(but correlation is 0.8 and knowing the value of x will give us the value of y).there is no statistical noise! it is just non_linear

#3:except the outlier,the relation between x and y is positive and it is linear but o utlier will cause lower correlation coefficient

#4:no relationship! knowing x does not tell anything about y ! also we have outlier

#Summary calculation (mean, sd), Pearson's correlation by set, and non-parametric, and p -value

#Correlation is a statistical measure that suggests the level of linear dependence between two variables

```
# summary_calc <- anscombe %>%
```

```
#   group_by(set) %>%
```

```
#   summarize(mean.x = mean(x), sd.x = sd(x), mean.y = mean(y), sd.y = sd(y), cor(x, y))
```

```
#
```

```
# summary_calc
```

```
summary_calc <- anscombe %>%
```

```
group_by(set) %>%
```

```
mutate(mean.x = mean(x), sd.x = sd(x), mean.y = mean(y), sd.y = sd(y))
```

```
summary_calc
```

```
## # A tibble: 44 x 7
## # Groups:   set [4]
##       x     y   set mean.x sd.x mean.y sd.y
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    10  8.04     1     9  3.20   7.50  1.96
## 2     8  6.95     1     9  3.20   7.50  1.96
## 3    13  7.58     1     9  3.20   7.50  1.96
## 4     9  8.81     1     9  3.20   7.50  1.96
## 5    11  8.33     1     9  3.20   7.50  1.96
## 6    14  9.96     1     9  3.20   7.50  1.96
## 7     6  7.24     1     9  3.20   7.50  1.96
## 8     4  4.26     1     9  3.20   7.50  1.96
## 9    12 10.8     1     9  3.20   7.50  1.96
## 10    7  4.82     1     9  3.20   7.50  1.96
## # ... with 34 more rows
```

```
#non-parametric, and p-value
ddply(anscombe, "set", summarise, corr=cor(x, y),
cor_spear = cor(x,y, method = "spearman"),
p.value = cor.test(x,y)$p.value)
```

```
##   set      corr cor_spear    p.value
## 1   1 0.8164205 0.8181818 0.002169629
## 2   2 0.8162365 0.6909091 0.002178816
## 3   3 0.8162867 0.9909091 0.002176305
## 4   4 0.8165214 0.5000000 0.002164602
```

#There are two main options for data aggregation:
#built-in functions, often referred to as the apply family of functions, the plyr add-on package

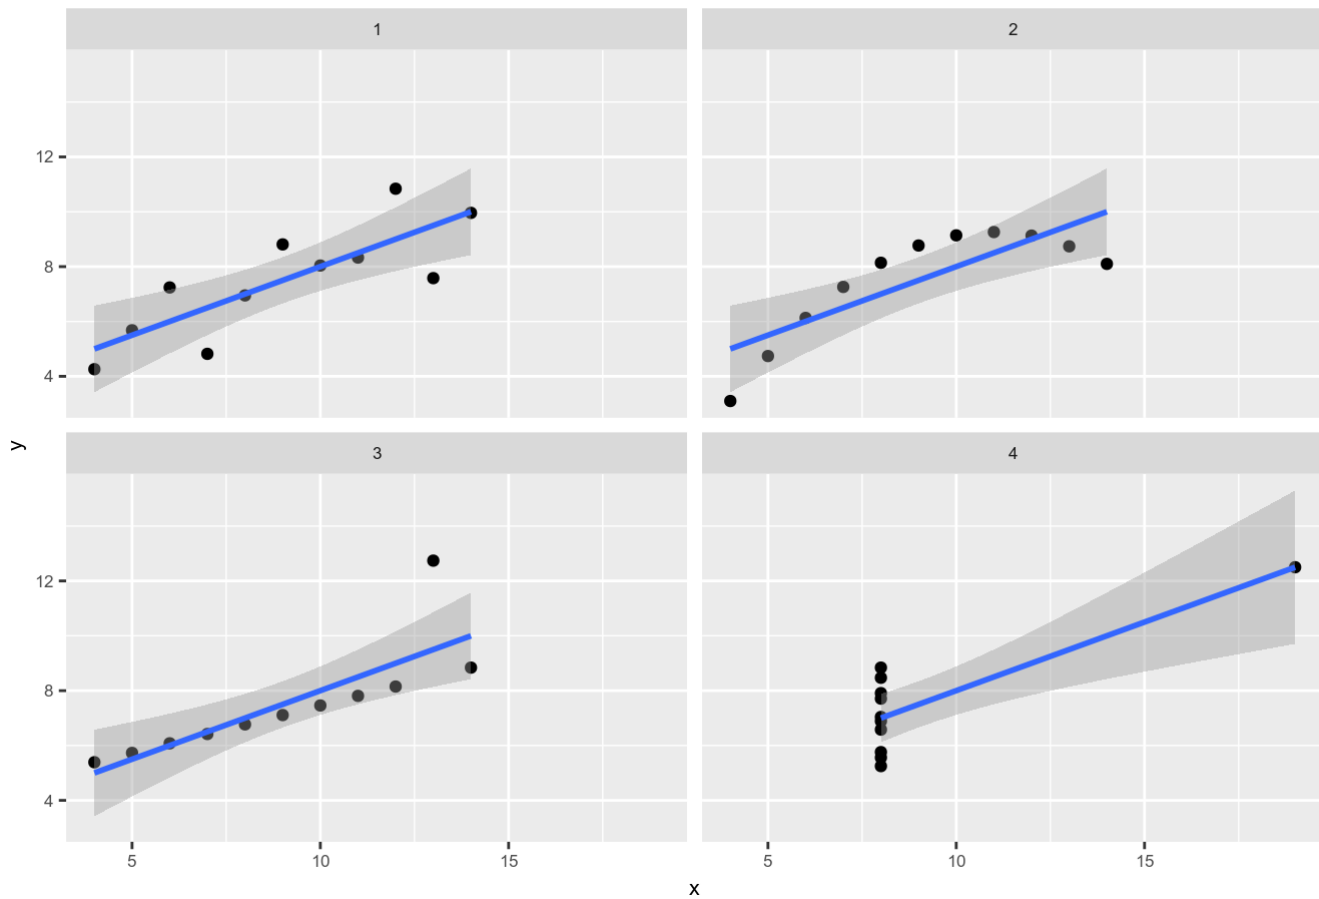
#The heart of plyr is a set of functions with names like this: XYply where X specifies what sort of input you're giving and Y specifies the sort of output you want.

```
# a = array, where matrices and vectors are important special cases
# d = data.frame
# l = list
```

The usage is very similar across these functions. Here are the main arguments:
.data is the first argument = the input
the next argument specifies how to split up the input into bits; it is does not exist when the input is a list, because the pieces are obviously the list components
then comes the function and further arguments needed to describe the computation to be applied to the bits
here ddply() will accept a data.frame, group it by "set", compute the correlation, then return the results as a data.frame.

```
q <- ggplot(anscombe, aes(x, y), color = as.factor(set)) +
  geom_point()+
  facet_wrap(~set)+
  theme(text = element_text(size = 8))+
  geom_smooth(method="lm")+
  ggtitle("Scatter plot for anscombe dataset ")
q
```

Scatter plot for anscombe dataset



#gray shading is standard error associating with the line!

#airquality data set

-200 is missing values in the data

```
airquality <- read.csv("/home/sedreh/Downloads/AirQualityUCI.csv", header=TRUE, sep=";", na.strings=c("-200", "-200,0"))
head(airquality)
```

```
##          Date      Time CO.GT. PT08.S1.CO. NMHC.GT. C6H6.GT. PT08.S2.NMHC.
## 1 10/03/2004 18.00.00    2,6      1360      150      11,9      1046
## 2 10/03/2004 19.00.00     2      1292      112      9,4      955
## 3 10/03/2004 20.00.00    2,2      1402      88      9,0      939
## 4 10/03/2004 21.00.00    2,2      1376      80      9,2      948
## 5 10/03/2004 22.00.00    1,6      1272      51      6,5      836
## 6 10/03/2004 23.00.00    1,2      1197      38      4,7      750
##   NOx.GT. PT08.S3.NOx. NO2.GT. PT08.S4.NO2. PT08.S5.O3.    T    RH    AH
## 1    166      1056      113      1692      1268 13,6 48,9 0,7578
## 2    103      1174      92      1559      972 13,3 47,7 0,7255
## 3    131      1140      114      1555      1074 11,9 54,0 0,7502
## 4    172      1092      122      1584      1203 11,0 60,0 0,7867
## 5    131      1205      116      1490      1110 11,2 59,6 0,7888
## 6     89      1337      96      1393      949 11,2 59,2 0,7848
##   X X.1
## 1 NA  NA
## 2 NA  NA
## 3 NA  NA
## 4 NA  NA
## 5 NA  NA
## 6 NA  NA
```

```
airquality <- data.frame(airquality)
```

```
# Drop the columns of the dataframe with NA value using select function in dplyr
airquality <- select (airquality,-c(X,X.1))
head(airquality)
```

```
##          Date      Time CO.GT. PT08.S1.CO. NMHC.GT. C6H6.GT. PT08.S2.NMHC.
## 1 10/03/2004 18.00.00    2,6      1360      150      11,9      1046
## 2 10/03/2004 19.00.00     2      1292      112      9,4      955
## 3 10/03/2004 20.00.00    2,2      1402      88      9,0      939
## 4 10/03/2004 21.00.00    2,2      1376      80      9,2      948
## 5 10/03/2004 22.00.00    1,6      1272      51      6,5      836
## 6 10/03/2004 23.00.00    1,2      1197      38      4,7      750
##   NOx.GT. PT08.S3.NOx. NO2.GT. PT08.S4.NO2. PT08.S5.O3.    T    RH    AH
## 1    166      1056      113      1692      1268 13,6 48,9 0,7578
## 2    103      1174      92      1559      972 13,3 47,7 0,7255
## 3    131      1140      114      1555      1074 11,9 54,0 0,7502
## 4    172      1092      122      1584      1203 11,0 60,0 0,7867
## 5    131      1205      116      1490      1110 11,2 59,6 0,7888
## 6     89      1337      96      1393      949 11,2 59,2 0,7848
```

```
# apply function anyNA() on all columns of airquality dataset and we can see that almost all columns (not Date,Time) have NA value!
```

```
str(airquality)
```

```
## 'data.frame':    9471 obs. of  15 variables:
## $ Date           : Factor w/ 392 levels "", "01/01/2005",...: 116 116 116 116 116 116
129 129 129 129 ...
## $ Time           : Factor w/ 25 levels "", "00.00.00",...: 20 21 22 23 24 25 2 3 4 5
...
## $ CO.GT.         : Factor w/ 103 levels "", "0,1", "0,2",...: 33 26 29 29 18 14 14 11
10 7 ...
## $ PT08.S1.CO.    : int   1360 1292 1402 1376 1272 1197 1185 1136 1094 1010 ...
## $ NMHC.GT.       : int   150 112 88 80 51 38 31 31 24 19 ...
## $ C6H6.GT.       : Factor w/ 408 levels "", "0,1", "0,2",...: 40 403 399 401 373 326 2
36 233 124 18 ...
## $ PT08.S2.NMHC.: int   1046 955 939 948 836 750 690 672 609 561 ...
## $ NOx.GT.        : int   166 103 131 172 131 89 62 62 45 NA ...
## $ PT08.S3.NOx.   : int   1056 1174 1140 1092 1205 1337 1462 1453 1579 1705 ...
## $ NO2.GT.        : int   113 92 114 122 116 96 77 76 60 NA ...
## $ PT08.S4.NO2.   : int   1692 1559 1555 1584 1490 1393 1333 1333 1276 1235 ...
## $ PT08.S5.O3.    : int   1268 972 1074 1203 1110 949 733 730 620 501 ...
## $ T              : Factor w/ 437 levels "", "-0,1", "-0,2",...: 67 64 50 41 43 43 44 3
8 38 34 ...
## $ RH             : Factor w/ 754 levels "", "10,0", "10,2",...: 376 364 427 487 483 47
9 455 487 484 489 ...
## $ AH             : Factor w/ 6684 levels "", "0,1847", "0,1862",...: 1897 1728 1854 20
57 2067 2046 1910 1963 1936 1864 ...
```

```
# here decimal is comma and comma is decimal
```

```
sub_clean_commas <- function(x) {gsub(",", ".", x)}
```

```
# T, RH, AH, C6H6.GT., CO.GT. these are factors due to commas
```

```
airquality$T <- sub_clean_commas(airquality$T)
airquality$RH <- sub_clean_commas(airquality$RH)
airquality$AH <- sub_clean_commas(airquality$AH)
airquality$C6H6.GT. <- sub_clean_commas(airquality$C6H6.GT.)
airquality$CO.GT. <- sub_clean_commas(airquality$CO.GT.)

str(airquality)
```

```
## 'data.frame':    9471 obs. of  15 variables:
## $ Date          : Factor w/ 392 levels "", "01/01/2005",...: 116 116 116 116 116 116
129 129 129 129 ...
## $ Time          : Factor w/ 25 levels "", "00.00.00",...: 20 21 22 23 24 25 2 3 4 5
...
## $ CO.GT.       : chr  "2.6" "2" "2.2" "2.2" ...
## $ PT08.S1.CO.  : int  1360 1292 1402 1376 1272 1197 1185 1136 1094 1010 ...
## $ NMHC.GT.     : int  150 112 88 80 51 38 31 31 24 19 ...
## $ C6H6.GT.     : chr  "11.9" "9.4" "9.0" "9.2" ...
## $ PT08.S2.NMHC.: int  1046 955 939 948 836 750 690 672 609 561 ...
## $ NOx.GT.      : int  166 103 131 172 131 89 62 62 45 NA ...
## $ PT08.S3.NOx. : int  1056 1174 1140 1092 1205 1337 1462 1453 1579 1705 ...
## $ NO2.GT.      : int  113 92 114 122 116 96 77 76 60 NA ...
## $ PT08.S4.NO2. : int  1692 1559 1555 1584 1490 1393 1333 1333 1276 1235 ...
## $ PT08.S5.03.  : int  1268 972 1074 1203 1110 949 733 730 620 501 ...
## $ T            : chr  "13.6" "13.3" "11.9" "11.0" ...
## $ RH           : chr  "48.9" "47.7" "54.0" "60.0" ...
## $ AH           : chr  "0.7578" "0.7255" "0.7502" "0.7867" ...
```

```
airquality <- airquality %>%
  select('C6H6.GT.', 'PT08.S1.CO.', 'PT08.S2.NMHC.', 'PT08.S3.NOx.', 'PT08.S4.NO2.', 'PT0
8.S5.03.', 'CO.GT.', 'NMHC.GT.', 'NOx.GT.', 'NO2.GT.', 'T', 'RH', 'AH')
head(airquality)
```

```
##   C6H6.GT. PT08.S1.CO. PT08.S2.NMHC. PT08.S3.NOx. PT08.S4.NO2. PT08.S5.03.
## 1    11.9      1360      1046      1056      1692      1268
## 2     9.4      1292       955      1174      1559       972
## 3     9.0      1402       939      1140      1555      1074
## 4     9.2      1376       948      1092      1584      1203
## 5     6.5      1272       836      1205      1490      1110
## 6     4.7      1197       750      1337      1393       949
##   CO.GT. NMHC.GT. NOx.GT. NO2.GT.    T    RH    AH
## 1    2.6    150    166    113 13.6 48.9 0.7578
## 2     2    112    103    92 13.3 47.7 0.7255
## 3    2.2    88    131    114 11.9 54.0 0.7502
## 4    2.2    80    172    122 11.0 60.0 0.7867
## 5    1.6    51    131    116 11.2 59.6 0.7888
## 6    1.2    38    89    96 11.2 59.2 0.7848
```

```
# convert all variables into numeric:
airquality <- data.matrix(airquality)
airquality <- data.frame(airquality)
str(airquality)
```

```
## 'data.frame':   9471 obs. of  13 variables:
## $ C6H6.GT.      : num  11.9 9.4 9 9.2 6.5 4.7 3.6 3.3 2.3 1.7 ...
## $ PT08.S1.CO.   : num  1360 1292 1402 1376 1272 ...
## $ PT08.S2.NMHC. : num  1046 955 939 948 836 ...
## $ PT08.S3.NOx.  : num  1056 1174 1140 1092 1205 ...
## $ PT08.S4.NO2.  : num  1692 1559 1555 1584 1490 ...
## $ PT08.S5.O3.   : num  1268 972 1074 1203 1110 ...
## $ CO.GT.        : num  2.6 2 2.2 2.2 1.6 1.2 1.2 1 0.9 0.6 ...
## $ NMHC.GT.      : num  150 112 88 80 51 38 31 31 24 19 ...
## $ NOx.GT.       : num  166 103 131 172 131 89 62 62 45 NA ...
## $ NO2.GT.       : num  113 92 114 122 116 96 77 76 60 NA ...
## $ T             : num  13.6 13.3 11.9 11 11.2 11.2 11.3 10.7 10.7 10.3 ...
## $ RH           : num  48.9 47.7 54 60 59.6 59.2 56.8 60 59.7 60.2 ...
## $ AH           : num  0.758 0.726 0.75 0.787 0.789 ...
```

```
# the data is ready for analysis now
```

```
sum(is.na(airquality))
```

```
## [1] 18183
```

```
airquality_summary <- colSums(is.na(airquality))
head(airquality_summary)
```

```
##      C6H6.GT.  PT08.S1.CO. PT08.S2.NMHC.  PT08.S3.NOx.  PT08.S4.NO2.
##           480           480           480           480           480
## PT08.S5.O3.
##           480
```

```
# We can see that NMHC.GT. column has too many NA's we must delete it, or else we will have little data to work with
airquality$NMHC.GT. <- NULL
airquality_clean <- na.omit(airquality)
head(airquality_clean)
```

```
##      C6H6.GT. PT08.S1.CO. PT08.S2.NMHC. PT08.S3.NOx. PT08.S4.NO2. PT08.S5.O3.
## 1      11.9      1360      1046      1056      1692      1268
## 2       9.4      1292      955      1174      1559      972
## 3       9.0      1402      939      1140      1555      1074
## 4       9.2      1376      948      1092      1584      1203
## 5       6.5      1272      836      1205      1490      1110
## 6       4.7      1197      750      1337      1393      949
##      CO.GT. NOx.GT. NO2.GT.      T      RH      AH
## 1      2.6     166     113 13.6 48.9 0.7578
## 2      2.0     103     92 13.3 47.7 0.7255
## 3      2.2     131    114 11.9 54.0 0.7502
## 4      2.2     172    122 11.0 60.0 0.7867
## 5      1.6     131    116 11.2 59.6 0.7888
## 6      1.2      89     96 11.2 59.2 0.7848
```

```
any(is.na(airquality_clean))
```

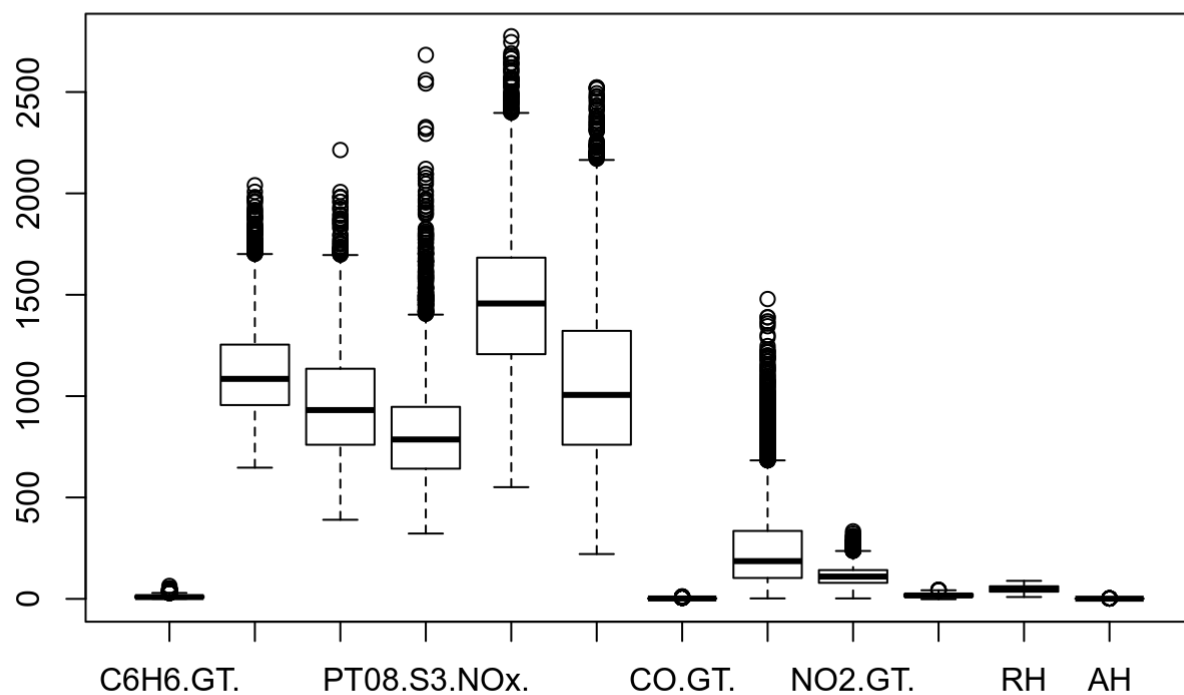


```
## [1] FALSE
```

```
#####Explore each variable independently#####
```


#Box plot:we should always perform an exploratory analysis of our variables before any formal modeling. This will give us a sense of our variable's distributions, any outliers, and any patterns that might be useful when constructing our eventual model.

```
plot <- boxplot(airquality_clean)
```



We can see that the data is out of scale

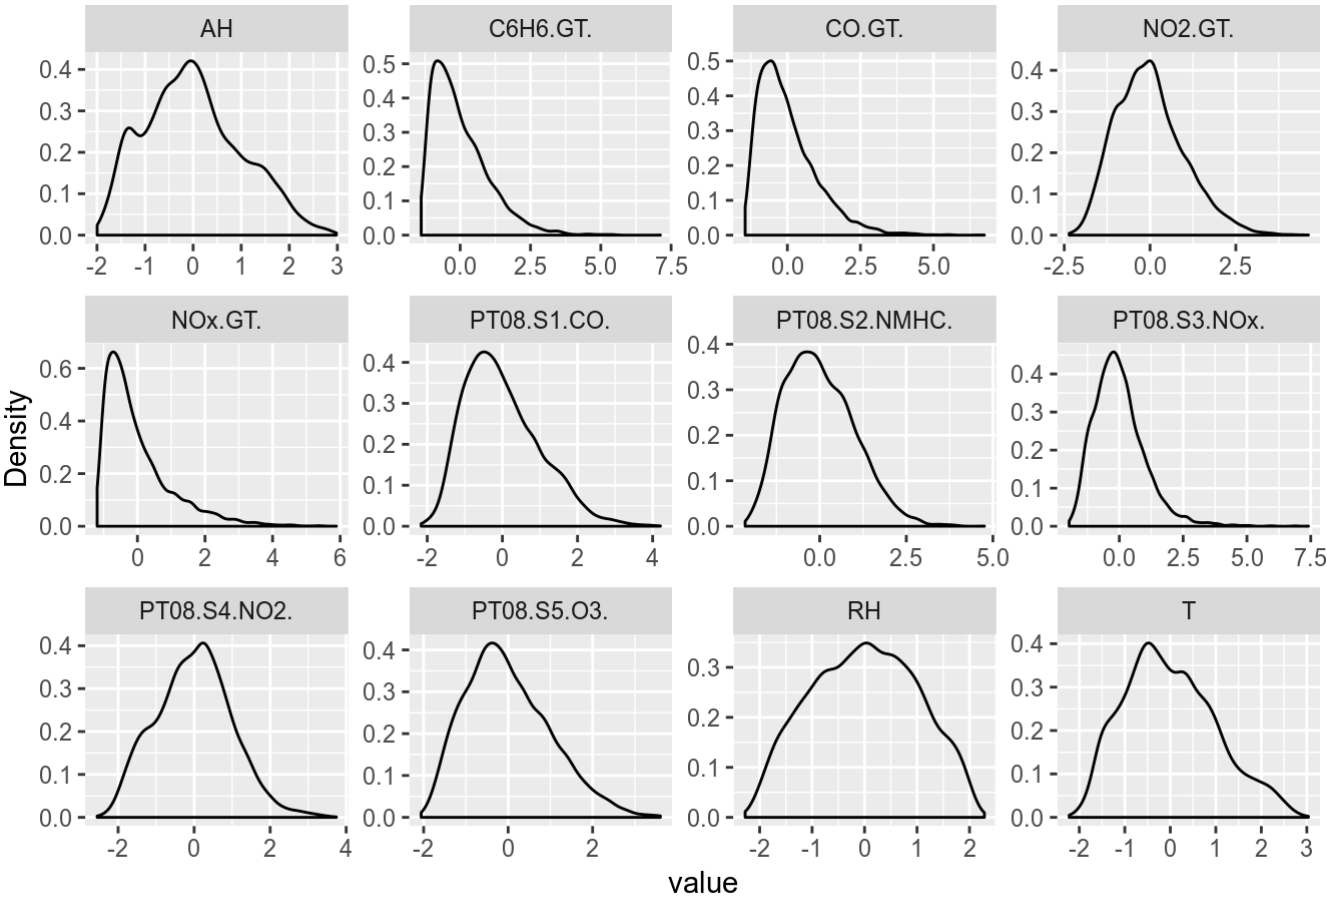
```
airquality_clean1 <- scale(airquality_clean, center = TRUE, scale = TRUE)
boxplot(airquality_clean1)
```




Some variable are in 10s but others are in 1000s this means that the variables in 10s will have no effect on the model even if we include/ exclude them, this is why we should bring them to the same scale to remove this effect.(normalization)

Now the data is in comparable scale.

#Density plot: To see the distribution of the predictor variable. Ideally, a close to normal distribution (a bell shaped curve), without being skewed to the left or right is preferred. Let see how to make each one of them.
`plot_density(airquality_clean1)`



```
#cross correlation
ggcorr(airquality_clean1, palette = "RdBu", label = TRUE)
```



```
#simple linear models with each predictor, check assumptions
#The summary for the linear model provides information regarding the quality of the model:

airquality_clean1 <- data.frame(airquality_clean1)

mod2 <- lm(C6H6.GT. ~ N02.GT., data=airquality_clean1)
summary(mod2)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ N02.GT., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3922 -0.4785 -0.0821  0.3463  5.7960
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.078e-15  9.574e-03   0.00      1
## N02.GT.      6.032e-01  9.574e-03  63.01 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7976 on 6939 degrees of freedom
## Multiple R-squared:  0.3639, Adjusted R-squared:  0.3638
## F-statistic: 3970 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod3 <-lm(C6H6.GT. ~ PT08.S1.CO., data=airquality_clean1)
summary(mod3)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S1.CO., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6014 -0.3027 -0.0245  0.2619  6.0196
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.947e-17  5.758e-03   0.0      1
## PT08.S1.CO.  8.774e-01  5.759e-03  152.4 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4797 on 6939 degrees of freedom
## Multiple R-squared:  0.7699, Adjusted R-squared:  0.7699
## F-statistic: 2.322e+04 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod4 <-lm(C6H6.GT. ~ PT08.S4.N02., data=airquality_clean1)
summary(mod4)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S4.N02., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.3317 -0.4794 -0.0616  0.3911  6.7388
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.111e-16  7.776e-03    0.00      1
## PT08.S4.N02.  7.618e-01  7.777e-03   97.96 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6479 on 6939 degrees of freedom
## Multiple R-squared:  0.5803, Adjusted R-squared:  0.5803
## F-statistic: 9596 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod5 <-lm(C6H6.GT. ~ PT08.S5.03., data=airquality_clean1)
summary(mod5)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S5.03., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1249 -0.3125 -0.0050  0.2829  4.4642
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5.385e-16  6.102e-03    0.0      1
## PT08.S5.03.  8.612e-01  6.103e-03   141.1 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5084 on 6939 degrees of freedom
## Multiple R-squared:  0.7416, Adjusted R-squared:  0.7415
## F-statistic: 1.991e+04 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod6 <-lm(C6H6.GT. ~ PT08.S3.N0x., data=airquality_clean1)
summary(mod6)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S3.NOx., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.2151 -0.4953 -0.1085  0.3089  5.9497
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.443e-16  8.259e-03    0.00      1
## PT08.S3.NOx. -7.257e-01  8.259e-03  -87.87 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.688 on 6939 degrees of freedom
## Multiple R-squared:  0.5267, Adjusted R-squared:  0.5266
## F-statistic: 7721 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod7 <-lm(C6H6.GT. ~ PT08.S2.NMHC., data=airquality_clean1)
summary(mod7)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S2.NMHC., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.15638 -0.12878 -0.06695  0.06782  2.44678
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -6.795e-16  2.223e-03    0.0      1
## PT08.S2.NMHC.  9.827e-01  2.223e-03  442.1 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1852 on 6939 degrees of freedom
## Multiple R-squared:  0.9657, Adjusted R-squared:  0.9657
## F-statistic: 1.954e+05 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod8 <-lm(C6H6.GT. ~ NOx.GT., data=airquality_clean1)
summary(mod8)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ NOx.GT., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.7014 -0.5321 -0.1491  0.3738  5.2240
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.374e-16  8.351e-03   0.00      1
## NOx.GT.      7.183e-01  8.352e-03  86.01 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6957 on 6939 degrees of freedom
## Multiple R-squared:  0.516, Adjusted R-squared:  0.5159
## F-statistic: 7398 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
mod9 <-lm(C6H6.GT. ~ CO.GT., data=airquality_clean1)
summary(mod9)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ CO.GT., data = airquality_clean1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0359 -0.2123 -0.0212  0.2005  6.5270
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.218e-16  4.412e-03   0.0      1
## CO.GT.      9.300e-01  4.412e-03  210.8 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3676 on 6939 degrees of freedom
## Multiple R-squared:  0.8649, Adjusted R-squared:  0.8649
## F-statistic: 4.443e+04 on 1 and 6939 DF, p-value: < 2.2e-16
```

```
#intercept = B0'
#slop = B1'
```

```
#Check all variables for their p-values

summary(mod2)$coefficients[4]
```

```
## [1] 0.009574462
```

```
summary(mod3)$coefficients[4]
```

```
## [1] 0.005758711
```

```
summary(mod4)$coefficients[4]
```

```
## [1] 0.007776722
```

```
summary(mod5)$coefficients[4]
```

```
## [1] 0.006102531
```

```
summary(mod6)$coefficients[4]
```

```
## [1] 0.008259098
```

```
summary(mod7)$coefficients[4]
```

```
## [1] 0.002223035
```

```
summary(mod8)$coefficients[4]
```

```
## [1] 0.008351533
```

```
summary(mod9)$coefficients[4]
```





```
## [1] 0.004412188
```

```
# options(digits=3)
# mod9[["coefficients"]][["(Intercept)"]] ==
# summary(mod9)$coefficients[4]
# # 3 is coefficients, 4 is value
# summary(mod9)$coefficients[4]
```

```
##### Diagnostic plots for the model #####
```

```
#Assumptions
#y-values or the errors are independent
#variation of observations around the regression line (residual SE) is constant(homoscedasticity)!
#for given value of x, y values (or the errors) are normally distributed

par(mfrow=c(2,2))
plot(mod2)
```

#x: predicted y values(y')

#y:errors or reseduals





#1)Residuals vs fitted values: if the linearity assumption is met we should see no pa tern here! the red line should be flat

#2) quantile quantile plot: y = observed and standardize residuals, x: theoritical re seduals.expected reseduals if errors/reseduals normally distributed.points shuld fall in diognal line!





#3)other 2 plots also can help to see non-linearities and non-constant varients and s o on!

#y' value for each data point!fitted values generate the residual! residual is the di fference between the actual observed value of the response variable and the expected value of the response according to our model!

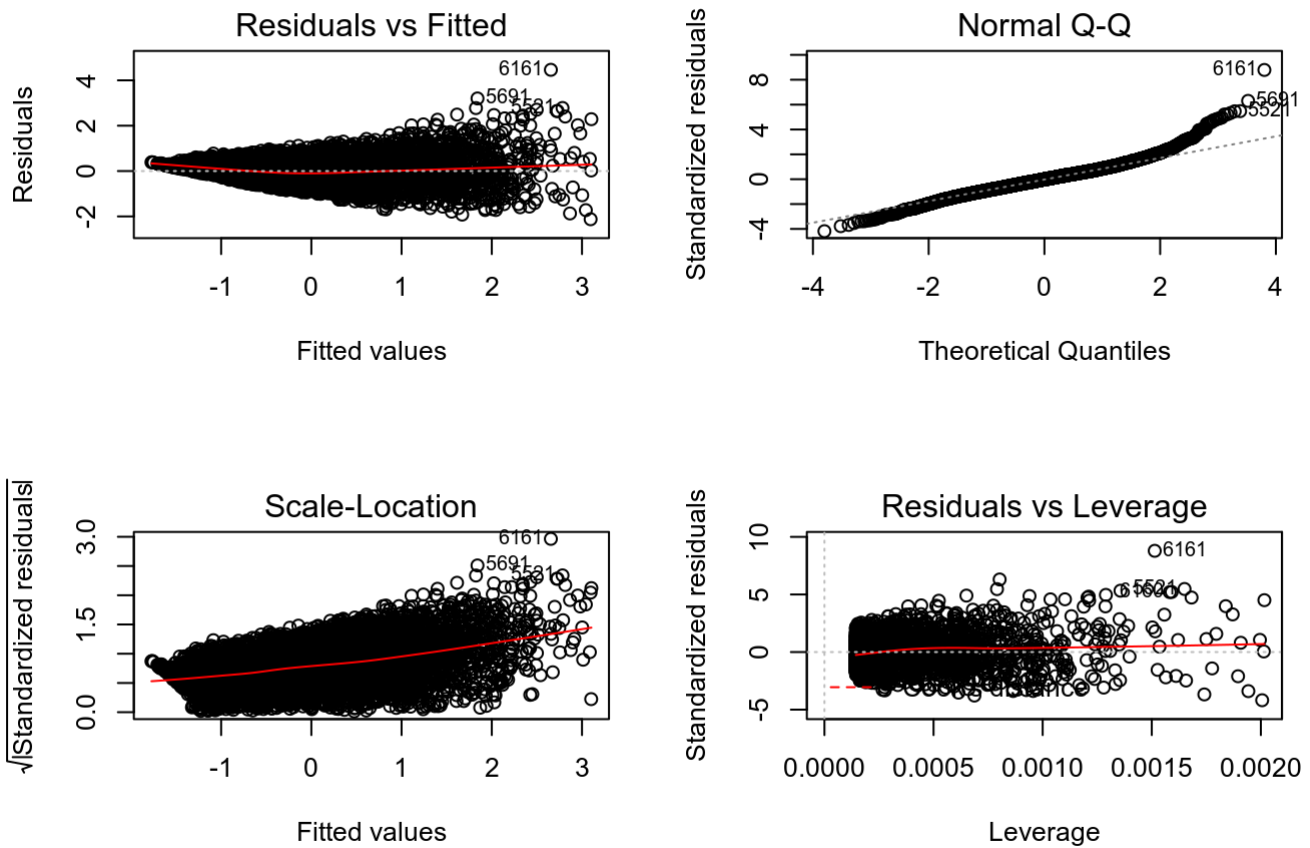
```
par(mfrow=c(2,2))
plot(mod3)
```



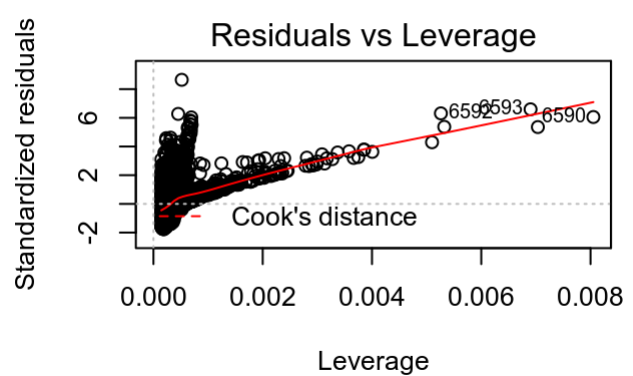



```
par(mfrow=c(2,2))
plot(mod4)
```



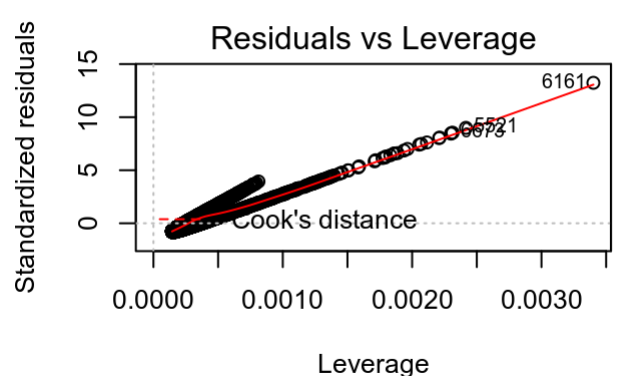



```
par(mfrow=c(2,2))
plot(mod5)
```







```
par(mfrow=c(2,2))
plot(mod6)
```







```
par(mfrow=c(2,2))
plot(mod7)
```



```
par(mfrow=c(2,2))
plot(mod8)
```



```
par(mfrow=c(2,2))
plot(mod9)
```



#For one of the models create train-test sets, plot the model, for the test set color real and predicted points differently; R^2 and p -value to title

TASK 11

```
library(caTools)
```

```
# let's use model 7
```

```
# splitting the dataset into test and train
```

```
split = sample.split(airquality_clean1, SplitRatio = 1/5) # 20 percent split
```

```
training_set = subset(airquality_clean1, split == TRUE)
```

```
test_set = subset(airquality_clean1, split == FALSE)
```

```
# Fitting and predicting on the training set which is 20% of the entire data
```

```
new_fit<- lm(C6H6.GT. ~ PT08.S2.NMHC., data=training_set)
```

```
summary(new_fit)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ PT08.S2.NMHC., data = training_set)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-0.15475	-0.12794	-0.07131	0.05908	2.43906

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.001335	0.005752	-0.232	0.817
PT08.S2.NMHC.	0.984610	0.005804	169.657	<2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1956 on 1155 degrees of freedom
## Multiple R-squared:  0.9614, Adjusted R-squared:  0.9614
## F-statistic: 2.878e+04 on 1 and 1155 DF, p-value: < 2.2e-16
```

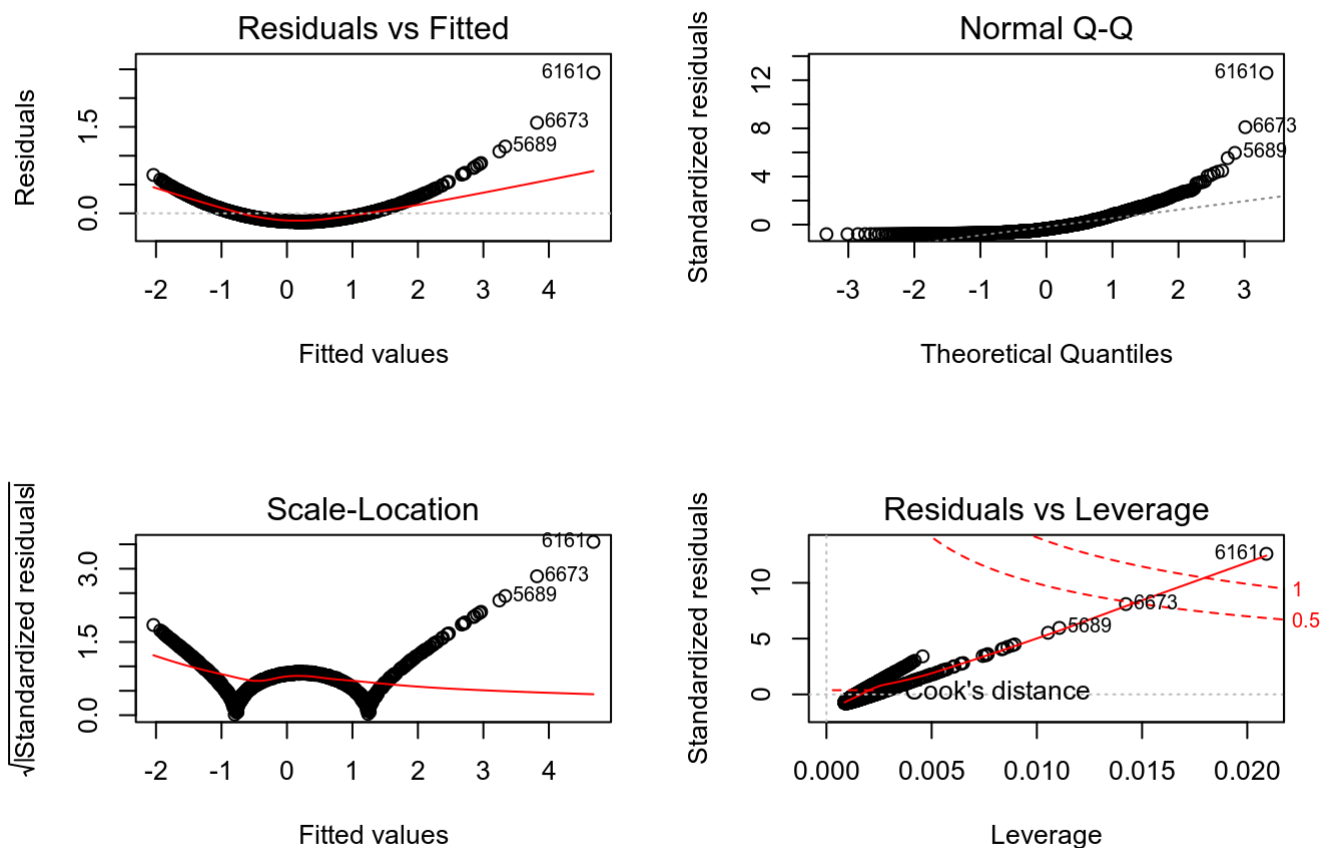
```
# data <- airquality_clean1
# sample <- sample.int(n = nrow(data), size = floor(.20*nrow(data)))
# train <- data[sample, ]
# test <- data[-sample, ]
# new_fit <- lm(C6H6.GT. ~ PT08.S2.NMHC., data=train)
# summary(new_fit)
```

```
# STATISTIC CRITERION
# R-Squared Higher the better (> 0.70)
# Adj R-Squared Higher the better
# F-Statistic Higher the better
# Std. Error Closer to zero the better
# t-statistic Should be greater 1.96 for p-value to be less than 0.05
# AIC Lower the better
# BIC Lower the better
# Mallows cp Should be close to the number of predictors in model
# MAPE (Mean absolute percentage error) Lower the better
# MSE (Mean squared error) Lower the better
# Min_Max Accuracy => mean(min(actual, predicted)/max(actual, predicted)) Higher the better
```

```
ggplot(data = training_set, aes(x =PT08.S2.NMHC. , y =C6H6.GT.))+
  geom_point()+
  geom_smooth(method = "lm")
```



```
# our model explains 96% variation in the data  
par(mfrow=c(2,2))  
plot(new_fit)
```

```
#### log transform this model
```

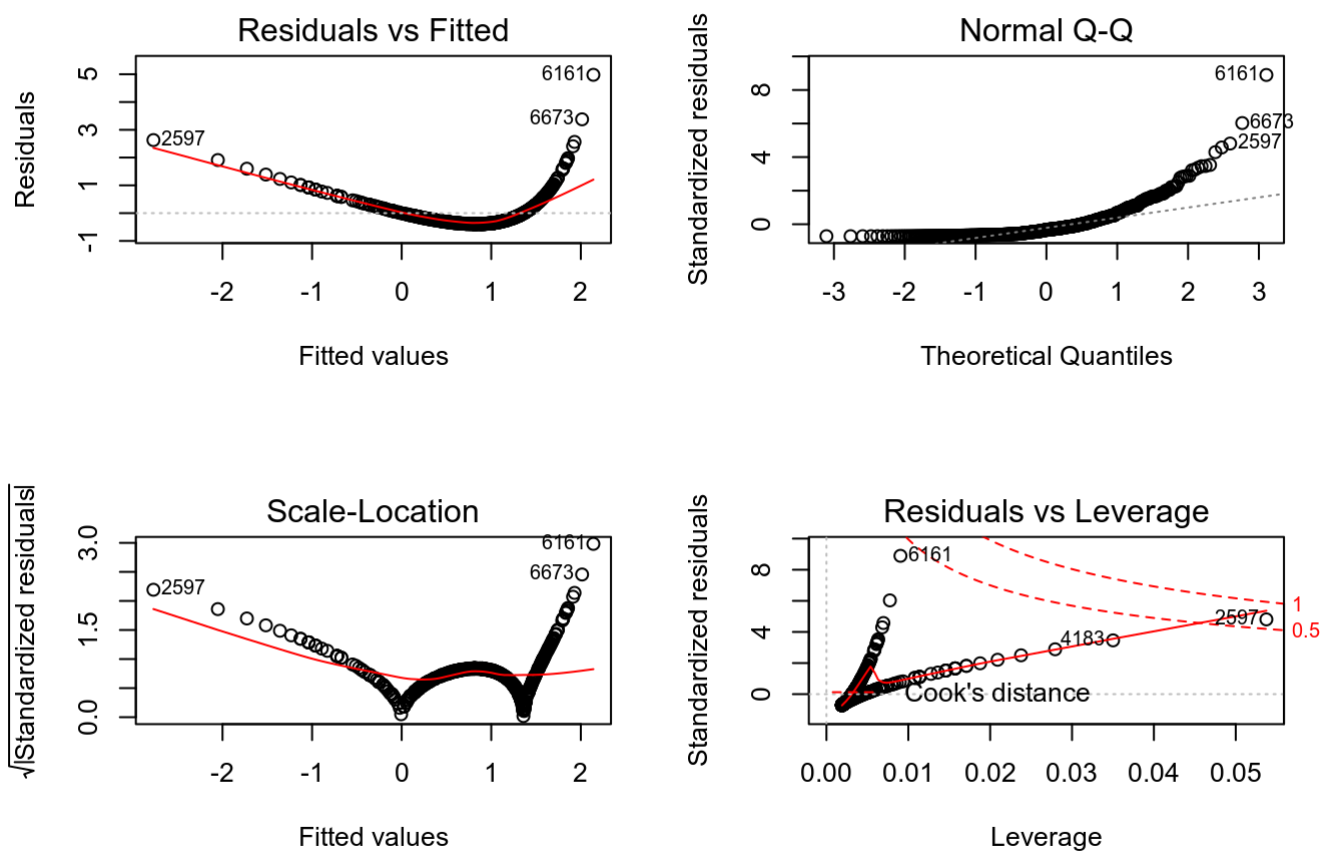
```
log_new_fit<- lm(C6H6.GT. ~ log(PT08.S2.NMHC.), data=training_set)
```

```
## Warning in log(PT08.S2.NMHC.): NaNs produced
```

```
summary(log_new_fit)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ log(PT08.S2.NMHC.), data = training_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.3996 -0.3490 -0.1857  0.1090  4.9789
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.17314    0.02771   42.33  <2e-16 ***
## log(PT08.S2.NMHC.) 0.62023    0.02217   27.97  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5623 on 525 degrees of freedom
## (630 observations deleted due to missingness)
## Multiple R-squared:  0.5985, Adjusted R-squared:  0.5977
## F-statistic: 782.5 on 1 and 525 DF, p-value: < 2.2e-16
```

```
# our model explains 96% variation in the data
par(mfrow=c(2,2))
plot(log_new_fit)
```



```
# log transforming data is not a good idea here
```

```
# pred <- predict(new_fit , new_data = test)
# head(pred)
#
# test$C6H6.GT._pred <- pred
# head(test)
```

```
## Let's Explore multicollinearity
```





```
adv_model <- lm(C6H6.GT. ~. -C6H6.GT., data = training_set)
summary(adv_model)
```

```
##
## Call:
## lm(formula = C6H6.GT. ~ . - C6H6.GT., data = training_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.86046 -0.09379 -0.01796  0.07165  2.52353
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.001872   0.004734   0.395   0.6926
## PT08.S1.CO.    0.009242   0.013747   0.672   0.5015
## PT08.S2.NMHC.  0.979525   0.025194  38.880 < 2e-16 ***
## PT08.S3.NOx.   0.108109   0.010581  10.217 < 2e-16 ***
## PT08.S4.NO2.   0.005746   0.020521   0.280   0.7795
## PT08.S5.O3.   -0.024121   0.013963  -1.728   0.0843 .
## CO.GT.         0.109609   0.015283   7.172 1.32e-12 ***
## NOx.GT.        0.074370   0.013666   5.442 6.44e-08 ***
## NO2.GT.       -0.070146   0.010219  -6.864 1.10e-11 ***
## T             -0.121615   0.017570  -6.922 7.41e-12 ***
## RH            -0.070749   0.013549  -5.222 2.10e-07 ***
## AH            0.072437   0.015530   4.664 3.46e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1609 on 1145 degrees of freedom
## Multiple R-squared:  0.9741, Adjusted R-squared:  0.9739
## F-statistic: 3919 on 11 and 1145 DF, p-value: < 2.2e-16
```

*# we can see that *** means 0.001 p values, these are the best predictors*
in our case we don't need to select good predictors as our model now explains 98% v
ariation, it is better than previous one
But if we have to choose some predictors, we apply a cut off say 0.05 and take all
the predictors that are less than this cut off, for example here we will select all
variable with three stars, remaining don't pass this cutoff

```
# PT08.S2.NMHC.  0.82449   0.02245  36.72 < 2e-16 ***
# PT08.S3.NOx.   0.07743   0.00822   9.43 < 2e-16 ***
# PT08.S4.NO2.   0.09561   0.01729   5.53 4.0e-08 ***
# CO.GT.         0.17500   0.01270  13.78 < 2e-16 ***
# NOx.GT.        0.09844   0.01067   9.23 < 2e-16 ***
# NO2.GT.       -0.07345   0.00835  -8.80 < 2e-16 ***
# T             -0.07541   0.01576  -4.78 1.9e-06 ***
# RH            -0.05864   0.01199  -4.89 1.1e-06 ***
```

```
par(mfrow=c(2,2))
plot(adv_model)
```



```
# let's make predictions from our model and compare with original data
```

```
predicted_values <- predict(adv_model, test_set)
```


```
original_values <- test_set$C6H6.GT.
```

```
total<- cbind(predicted_values, original_values)
```

```
colnames(total) <- c('Predicted by model', 'Original')
```

```
# plot
```

```
plot_density(total)
```



we can see that the distribution plots are similar for the model predicted values and original values, our model is thus effective

```
# library(corrplot)
# visual_cor <- function(d) {
#   cormat <- cor(d , use = "pairwise.complete.obs")
#   pvalmat <- cor.mtest(d)$p
#
#   corrplot(abs(cormat),
#             method = 'color',
#             order = 'hslust',
#             addCoef.col = "black",
#             tl.col = "black", tl.srt = 45,
#             pmat = pvalmat,
#             sig.level = 0.05,
#             insig = "blank",
#             diag = "FALSE")
# }
# visual_cor(training_set[, -c(1:4)])
```

#####usefulexplanation#####

#before creating model do some exploratory data analysis!

#1)looking at data

#2)visualization

#3)summary statistics

#the role of modeling is to explore the relationship between variables. using scatter plot you can show this relationship.

#Build simple linear models with each predictor. Linear regression is used to predict the value of an outcome variable Y based on one or more input predictor variables X . The aim is to establish a linear relationship (a mathematical formula) between the predictor variable(s) and the response variable, so that, we can use this formula to estimate the value of the response Y , when only the predictors (X s) values are known.

#the aim of linear regression is to model a continuous variable Y as a mathematical function of one or more X variable(s), so that we can use this regression model to predict the Y when only the X is known.

$Y = \beta_1 + \beta_2 X + \epsilon$

#where, β_1 is the intercept and β_2 is the slope. Collectively, they are called regression coefficients. ϵ is the error term, the part of Y the regression model is unable to explain.

#consists of 1231 observations(rows) and 15 variables (columns)

#Before we begin building the regression model, it is a good practice to analyze and understand the variables. The graphical analysis and correlation study below will help with this.

#the aim of this exercise is to build a simple regression model that we can use to predict NMHC.GT. by establishing a statistically significant linear relationship with all other variables. But before jumping in to the syntax, let's try to understand these variables graphically. Typically, for each of the independent variables (predictors), the following plots are drawn to visualize the following behavior:

#It adds a small amount of random variation to the location of each point, and is a useful way of handling overplotting caused by discreteness in smaller datasets.

#we can understand the correlation between variables(positive or negative) using correlation coefficient(summary statistic between negative one and one measuring strength linear association of 2 numerical variables(for example in line that as well as x increase y will decrease we can tell that it is negatively line slope!) when x and y behave independently we can say that there is 0 relationship)

#cor() function takes 2 numerical variable and gives correlation (can not be greater than 1):data %>% summarize(correlation = cor(x,y)) and when the value is negative we can say that it is negative relationship(weakly negative). correlation near the 1 will show positive strong relationship

#prediction with modeling: if i give input x to $f()$ can i get prediction of y (hat) that is close to the true value of y !

#both variables are numeric.the name of output will be response variable(and it is the quantity that we believe might be related to the input!)

#explanatory variable:something that you think might be related to the response

x is independent(predictor)

y is dependent (response)

#graphical representation: response= y axis, explanatory/predictor= x axis

#The relationship between two variables may not be linear. In these cases we can sometimes see strange and even inscrutable patterns in a scatterplot of the data. Sometimes there really is no meaningful relationship between the two variables. Other times, a careful transformation of one or both of the variables can reveal a clear relationship.

#we can understand that how outliers can affect the results of a linear regression model and how we can deal with them. For now, it is enough to simply identify them and note how the relationship between two variables may change as a result of removing outliers.

#we can summarize relationship between 2 variable with fitting line(linear regression line that separating signals from the noise) by adding `geom_smooth()` with `method= lm` for linear model and `se = FALSE` to meet standard error bars to the ggplot.

#to show the relation between x and y we need to draw a line! and understand which line will fit better! we need a numerical measurement of how good fit of each possible line is? in regression we use "least square criterion" to find the best fit line! line that minimize the sum square of distances between the line and a set of data points! a unique line exist! that line is called "least square regression line". we can add the line to the plot using the `geom_smooth` function!

#regression model combine some explanatory variables to estimate single numerical response variable!

#The function used for building linear models is `lm()`. The `lm()` function takes in two main arguments, namely: 1. Formula 2. Data.

#making model for prediction! like making model for prediction of price of apartment based on some features like size, condition, number of floors and so on!

#slop: شیب خط رگرشن

#response = f(explanatory) + noise

#regression = intercept + f(slop*explanatory)+noise <- $y = B_0 + B_1.x + e$

$y' = B_0' + B_1'.X$ y = actual observed values of the response, y' = expected values of the response based on the model

#residuals : difference between y and y' $e' = y - y'$

e is unknown quantity(noise) but e' is known estimate of e

#residuals: is the difference between the observed y value and predicted or fitted y value (y')

#standard deviation of this errors or residuals called residuals standard error in R and is presented in the model summary. residuals or error terms are useful for checking the model assumptions

#assumptions:

#y-values or the errors are independent

#y-values can be expressed as a linear function of the x-values

#variation of observations around the regression line (residual SE) is constant!

#for given value of x, y values are normally distributed

#Null hypothesis is a general statement that there is no relationship between two measured phenomena.